Introduction to Kubernetes

Ihor Dvoretskyi, Developer Advocate Cloud Native Computing Foundation @idvoretskyi



LinuxFoundationX: LFS158x



- This is a brief overview of Kubernetes, based on the Linux Foundation Course LFS158x
- Full course is available on <u>EDX</u>
 <u>for free</u>





Containers



What Are Containers?



app bin/lib

Desktop

Dev VM

QA Env.

Public Cloud

Private Cloud

Customer Site



Why Containers?



- Portability
- Immutability
- Ease of update (layers)
- Pre-packaged containers
- Lightweight (create, destroy)
- Excellent for microservices
- Standard repositories



Containers vs Virtual Machines









Container Management

数

- Containers are powerful, running one container is useless
- Applications using microservices will include components, bundled into multiple containers
- To manage them, the Container scheduling & orchestration tools are required



Kubernetes



Docker Swarm/ Compose



Kubernetes



What Is Kubernetes?

- Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications.
- Originally developed by Google (inspired by the internal system called Borg)





Key Kubernetes Features

- Self-healing
- Horizontal scaling
- Service discovery and Load balancing
- Automated rollouts and rollbacks
- Secrets and configuration management
- Storage orchestration

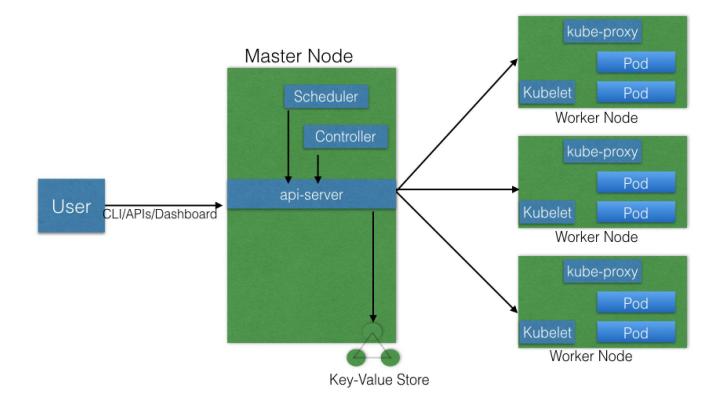




Kubernetes Architecture



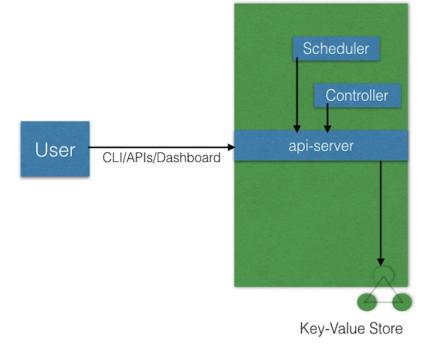
Kubernetes Architecture





Master Node aka Control Plane

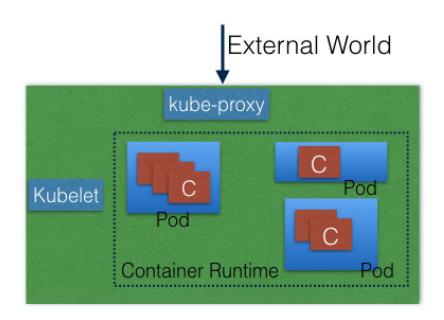
- The master node is responsible for managing the Kubernetes cluster, and it is the entry point for all administrative tasks
- User can communicate to the Master Node via the CLI, the GUI (Dashboard), or via APIs.
- For fault tolerance purposes, there can be more than one master node in the cluster (Multi-Master).





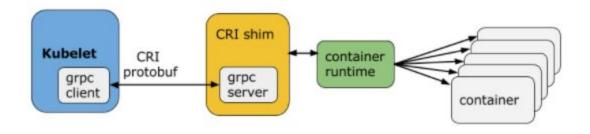
Worker Node

- A worker node is a machine (VM, physical server, etc.) which runs the applications using Pods and is controlled by the master node
- Pods are scheduled on the worker nodes, which have the necessary tools to run and connect them
- A Pod is the atomic of the Kubernetes workloads





Container Runtime Interface





Kubernetes Configuration



Kubernetes Configurations



- All-in-One Single-Node Installation
- Single-Node etcd, Single-Master, and Multi-Worker Installation
- Single-Node etcd, Multi-Master, and Multi-Worker Installation
- Multi-Node etcd, Multi-Master, and Multi-Worker Installation



Localhost Installation



- There are a few localhost installation options available to deploy single- or multi-node Kubernetes clusters on our workstation/laptop:
 - Minikube
 - <u>Ubuntu on LXD</u>
- Minikube is the preferred and community-recommended way to create an all-in-one Kubernetes setup



Non-Local Installations



- On-Premise Installation
 - o On-Premise VMs
 - o On-Premise Bare Metal
- Public Cloud Installation
 - Managed Solutions
 - Google Kubernetes Engine (GKE)
 - Azure Container Service (AKS)
 - Amazon Elastic Container Service for Kubernetes (EKS)
 - Self-managed Solutions
 - Kubeadm
 - Kops
 - Kubespray



Minikube



Minikube

- The community-recommended way to deploy Kubernetes locally
- Minikube runs inside a VM on:
 - Windows
 - Mac
 - Linux
 - Also available without VM, being run directly on host (Docker is required)



Installing Minikube (Requirements)



- kubectl
- macOS
 - Hyperkit driver, xhyve driver, VirtualBox, or VMware Fusion
- Linux
 - VirtualBox or KVM
 - NOTE: Minikube also supports a --vm-driver=none option that runs the Kubernetes components on the host and not in a VM
- Windows
 - VirtualBox or Hyper-V
- VT-x/AMD-v virtualization must be enabled in BIOS
- Internet connection on first run



Installing Minikube



- Use the link to find installation instructions for your OS: https://github.com/kubernetes/minikube#installation
- Linux:
 - curl -Lo minikube
 https://storage.googleapis.com/minikube/releases/latest/minikube-linux-amd64 && chmod +x
 minikube && sudo cp minikube /usr/local/bin/ && rm minikube
- macOS (with brew):
 - brew cask install minikube
- Windows (with chocolatey):
 - choco install minikube



Accessing Minikube



Any healthy running Kubernetes cluster can be accessed via one of the following methods:

- Command Line Interface (CLI)
- Graphical User Interface (GUI)
- APIs



CLI: kubectl



- Use the link to find installation instructions for your OS:
 - https://kubernetes.io/docs/tasks/tools/install-kubectl/
- Check if it is working:
 - kubectl help



UI: dashboard



- Enabled by default in Minikube, accessible via:
 - o minikube dashboard





Q Search

+ CREATE



Cluster

Namespaces

Nodes

Persistent Volumes

Roles

Storage Classes

Namespace

default ▼

Overview

Workloads

Cron Jobs

Daemon Sets

Deployments

Jobs

Pods

Replica Sets

Replication Controllers

Discovery and Load Balancing



Config and Storage

Secrets			Ŧ
Name 🕏	Туре	Age 🔷	
default-token-pdf9q	kubernetes.io/service-account-token	7 minutes	•



Using the 'kubectl proxy' Command



\$ kubectl proxy

Starting to serve on 127.0.0.1:8001

\$ curl http://localhost:8001/



Kubernetes Objects



Kubernetes Objects



Kubernetes has a very rich object model, with which it represents different persistent entities in the Kubernetes cluster. Those entities describe:

- What containerized applications we are running and on which node
- Application resource consumption
- Different policies attached to applications, like restart/upgrade policies, fault tolerance, etc.

Kubernetes objects include:

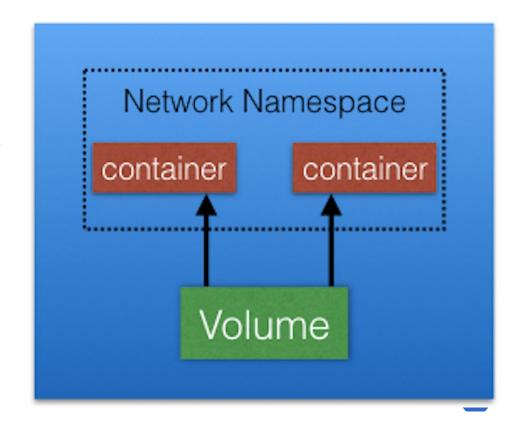
- Pods
- ReplicaSets
- Deployments
- Namespaces



Pods

A Pod is the smallest and simplest Kubernetes object. It is the unit of deployment in Kubernetes, which represents a single instance of the application. A Pod is a logical collection of one or more containers, which:

- Are scheduled together on the same host
- Share the same network namespace
- Mount the same external storage (volumes).



Labels

- Labels are key-value pairs that can be attached to any Kubernetes objects (e.g. Pods).
- Labels are used to organize and select a subset of objects, based on the requirements in place.
- Many objects can have the same Label(s).
- Labels do not provide uniqueness to objects.







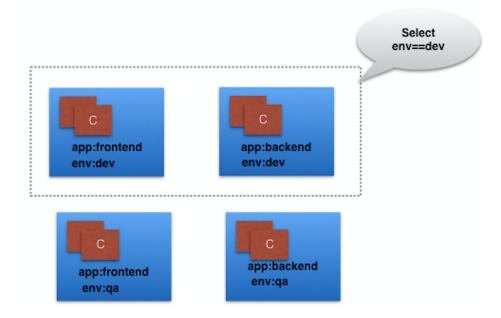




Label Selectors

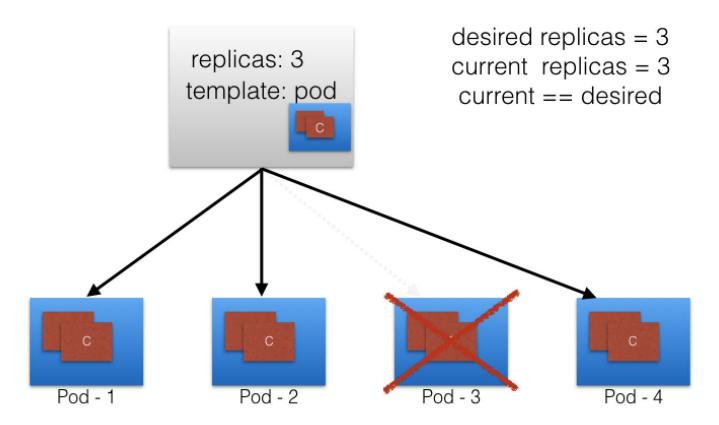
With Label Selectors, we can select a subset of objects. Kubernetes supports two types of Selectors:

- Equality-Based Selectors
 - Equality-Based Selectors allow filtering of objects based on Label keys and values
 - env==dev
- Set-Based Selectors
 - Set-Based Selectors allow filtering of objects based on a set of values
 - env in (dev,qa)





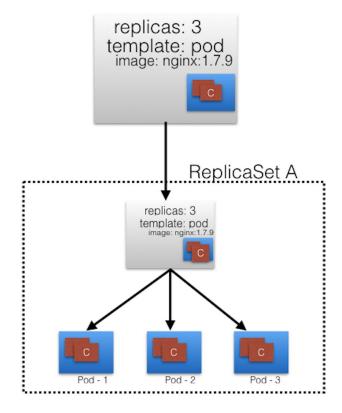
ReplicaSets





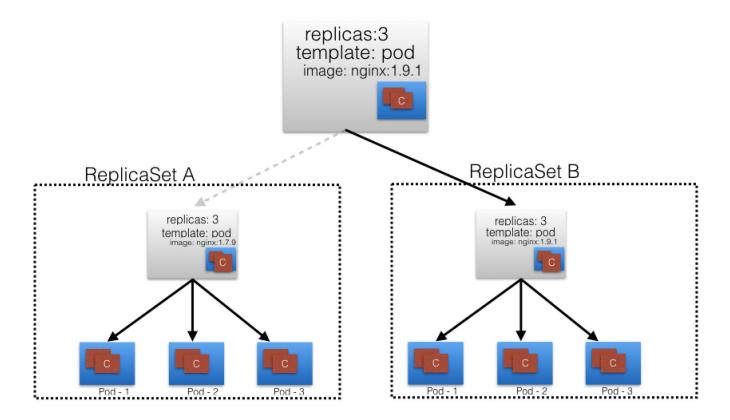
Deployments

- Deployment objects provide declarative updates to Pods and ReplicaSets
- The DeploymentController is part of the master node's controller manager, and it makes sure that the current state always matches the desired state.





Deployment rollout





Namespaces

- If we have numerous users
 whom we would like to organize
 into teams/projects, we can
 partition the Kubernetes cluster
 into sub-clusters using
 Namespaces.
- Kubernetes creates two default Namespaces:
 - kube-system
 - default

\$ kubectl get namespaces

NAME STATUS AGE

default Active 11h

kube-public Active 11h

kube-system Active 11h



Authentication, Authorization, and Admission Control



Overview

- Authentication
 - Logs in a user.
- Authorization
 - Authorizes the API requests added by the logged-in user.
- Admission Control
 - Software modules that can modify or reject the requests based on some additional checks, like Quota.



Authentication

Kubernetes has two kinds of users (while they are not *users* in a common meaning):

Normal Users

• They are managed outside of the Kubernetes cluster via independent services like User/Client Certificates, a file listing usernames/passwords, Google accounts, etc.

Service Accounts

- With Service Account users, in-cluster processes communicate with the API server to perform different operations
- Most of the Service Account users are created automatically via the API server, but they can also be created manually.



Authorization

- After a successful authentication, users can send the API requests to perform different operations
- Then, those API requests get authorized by Kubernetes using various authorization modules.



Authorization modules

Node Authorizer

 Node authorization is a special-purpose authorization mode which specifically authorizes API requests made by kubelets.

Attribute-Based Access Control (ABAC) Authorizer

 With the ABAC authorizer, Kubernetes grants access to API requests, which combine policies with attributes.

Webhook Authorizer

• With the Webhook authorizer, Kubernetes can offer authorization decisions to some third-party services, which would return true for successful authorization, and false for failure.

Role-Based Access Control (RBAC) Authorizer

 In Kubernetes, we can have different roles that can be attached to subjects like users, service accounts, etc.



Admission Control

- Admission control is used to specify granular access control policies, which include allowing privileged containers, checking on resource quota, etc.
- We force these policies using different admission controllers, like ResourceQuota, AlwaysAdmit, DefaultStorageClass, etc.
- They come into effect only after API requests are authenticated and authorized.



Services

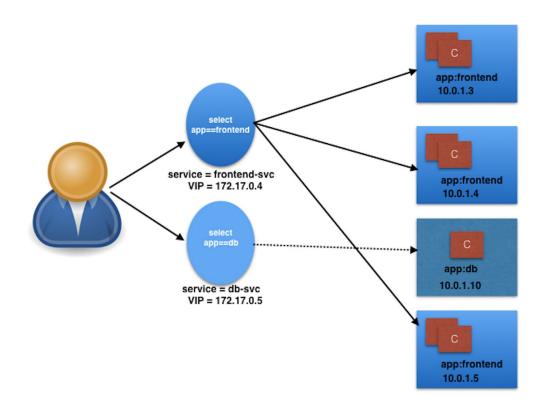


Overview

- To access the application, a user/client needs to connect to the Pods
- As Pods are ephemeral in nature, resources like IP addresses allocated to it cannot be static
- Pods could die abruptly or be rescheduled based on existing requirements
- Unexpectedly, the Pod to which the user/client is connected dies, and a new Pod is created by the controller
- The new Pod will have a new IP address, which will not be known automatically to the user/client of the earlier Pod
- Kubernetes provides a higher-level abstraction called **Service**, which logically groups Pods and a policy to access them.



Grouping of Pods using the Service object



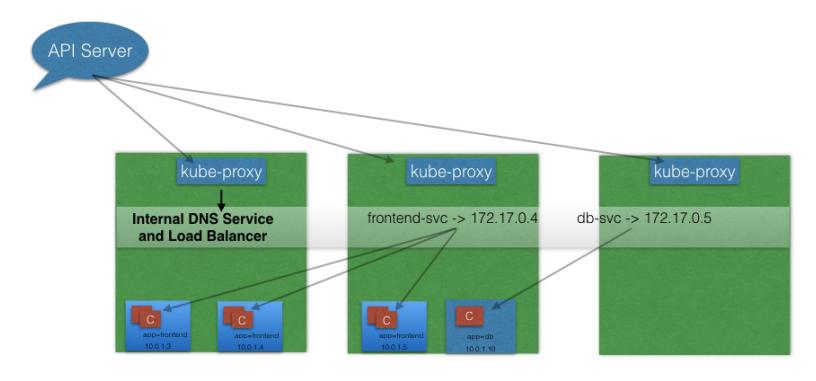


Kube-proxy

- All of the worker nodes run a daemon called kube-proxy, which watches the API server on the master node for the addition and removal of Services and endpoints.
- For each new Service, on each node, kube-proxy configures the iptables rules to capture the traffic for its ClusterIP and forwards it to one of the endpoints.
- When the service is removed, kube-proxy removes the iptables rules on all nodes as well.



kube-proxy, Services, and Endpoints





ServiceType

- While defining a Service, we can also choose its access scope. We can decide whether the Service:
 - Is only accessible within the cluster
 - Is accessible from within the cluster and the external world
 - Maps to an external entity which resides outside the cluster.
- Access scope is decided by ServiceType, which can be mentioned when creating the Service.

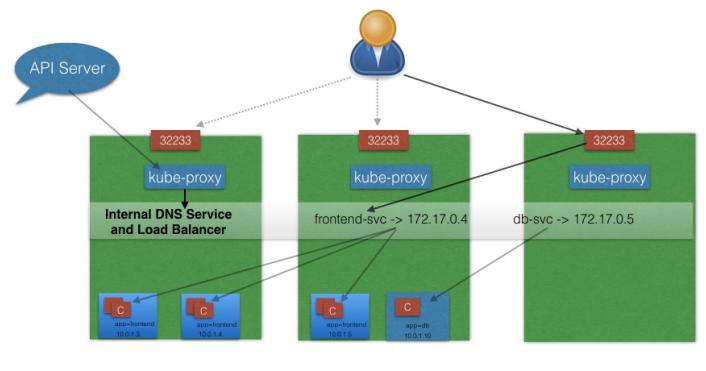


ServiceType: ClusterIP and NodePort

- ClusterIP is the default ServiceType.
 - A Service gets its Virtual IP address using the ClusterIP
 - That IP address is used for communicating with the Service and is accessible only within the cluster.
- The NodePort ServiceType is useful when we want to make our Services accessible from the external world.
 - The end-user connects to the worker nodes on the specified port, which forwards the traffic to the applications running inside the cluster.
 - To access the application from the external world, administrators can configure a reverse proxy outside the Kubernetes cluster and map the specific endpoint to the respective port on the worker nodes.



ServiceType: ClusterIP and NodePort



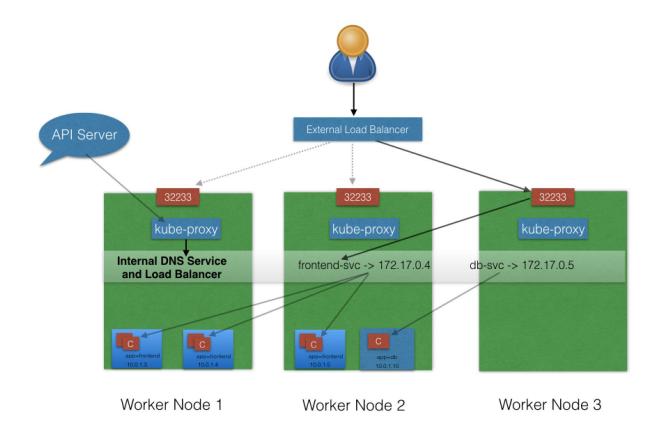


ServiceType: LoadBalancer

- With the LoadBalancer ServiceType:
 - NodePort and ClusterIP Services are automatically created, and the external load balancer will route to them
 - The Services are exposed at a static port on each worker node
 - The Service is exposed externally using the underlying cloud provider's load balancer feature.
- The LoadBalancer ServiceType will only work if the underlying infrastructure supports the automatic creation of Load Balancers and have the respective support in Kubernetes. Examples:
 - Google Cloud Platform
 - > AWS



ServiceType: LoadBalancer



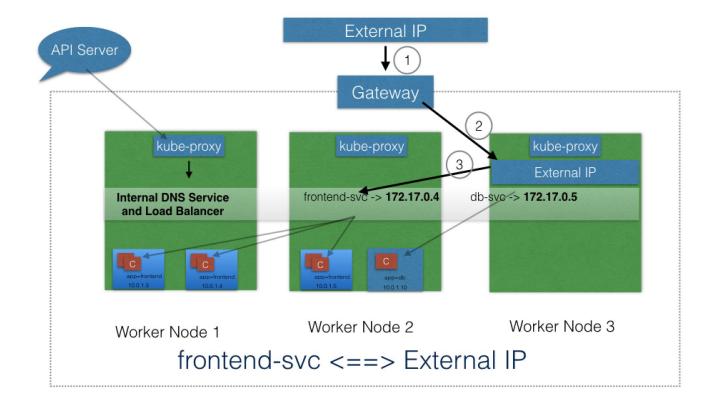


ServiceType: ExternalIP

- A Service can be mapped to an ExternalIP address if it can route to one or more of the worker nodes.
- Traffic that is ingressed into the cluster with the ExternalIP (as destination IP)
 on the Service port, gets routed to one of the Service endpoints.



ServiceType: ExternalIP





Deploying a Standalone Application



Deploying the Application Using the CLI

- Use the sample YAML file available at:
 - https://github.com/idvoretskyi/intro-to-kubernetes/blob/master/app/webserver.yaml
- Creating a webserver Deployment:
 - kubectl create -f webserver.yaml
- This will also create a ReplicaSet and Pods, as defined:
 - kubectl get replicasets



Creating a Service and Exposing It

- Use the sample YAML file available at:
 - https://github.com/idvoretskyi/intro-to-kubernetes/blob/master/app/webserver.yaml
- Using kubectl, create the Service:
 - kubectl create -f webserver-svc.yaml
- List the Services:
 - kubectl get svc
- To get more details about the service:
 - kubectl describe svc web-service
- Test the app:
 - minikube service web-service --url #to get an IP address of the service
 - o curl http://192.168.99.100:32742



Kubernetes Community



Getting Started with the Community



- Weekly Meetings
- Meetup Groups
- Slack Channels
- Mailing Lists
- Special Interest Groups (SIGs)
- Stack Overflow
- CNCF Events
 - KubeCon + CloudNativeCon Europe
 - KubeCon + CloudNativeCon China
 - KubeCon + CloudNativeCon North America



Next



- Expand your Kubernetes knowledge and skills by enrolling in paid courses offered by The Linux Foundation:
 - self-paced LFS258 <u>Kubernetes Fundamentals</u>
 - instructor-led LFS458 <u>Kubernetes Administration</u>
- Prepare for:
 - <u>Certified Kubernetes Administrator exam (CKA)</u>
 - <u>Certified Kubernetes Application Developer Program (CKAD)</u>, offered by the <u>Cloud Native Computing Foundation</u>



Questions?

@idvoretskyi kubernetes.io cncf.io

